

Microprocessor History

- Mid 1970
 - Microprocessor introduced.
- Mid 1970 – 1980s
 - Two factors influenced architecture:
 - Microprogramming and Complex Instructions
 - Ferrite memory core had a long access time.
 - The slow main store held complex instructions.
 - Fetching and executing microprograms from the much faster microprogram memory within the CPU was advantageous.
 - The complex instructions were seen to help programmers.
- Today, most of the advantages of microprogramming have evaporated due to the low access time of system memory and cache systems.

The RISC Revolution

- 1980s
 - Initial reaction against the trend towards complex instructions:
 - IBM's 801 architecture
 - In 1974 John Cocke (IBM) starting working on RISC like architectures.
 - Berkeley: David Patterson and David Ditzel
 - Coined RISC

RISCy Research

- Research conducted in the late 1970s by Fairclough demonstrated that the relative frequency with which different classes of instructions are executed is not uniform and some types of instructions are executed much more frequently than others.

Instruction Type	Frequency
Data Movement	45.28
Instruction Flow Modification (Branch, Call, Return)	28.73
Arithmetic	10.75
Compare	5.92
Logical	3.91
Shift	2.93
Bit Manipulation	2.05
IO and miscellaneous	0.44

- Conclusion: optimize the highly frequent instructions.

RISCy Research

- Tanenbaum reported that 56% of all constant values lie in the range from -15 to +15, and 98% of all constant values lie in the range from -511 to +511.
 - Optimum size of an instruction.
- Other research showed that 12 words of storage are sufficient for parameter passing to and from subroutines:
 - 12 internal registers dedicated for subroutine parameter passing.

Characteristics of the RISC Architecture

1. Have sufficient on-chip registers to overcome processor-memory bottleneck.
2. Three address, register to register instruction set architecture.
`OPERATION Ra, Rb, Rc`
1. Facilitate passing of parameters to and from subroutines
 - For example, internal registers.
1. Program flow instructions are implemented efficiently.

Characteristics of the RISC Architecture

5. Don't implement infrequently used instructions.
 - Complex instructions waste hardware real estate.
 - Complex instructions increase design, fabrication, and test times.

6. Execute an instruction in a single clock cycle, through
 - Regularity in the instruction set.
 - Fixed size instruction.
 - All instructions take the same number of clock cycles to execute.
 - Pipelining.

Characteristics of the RISC Architecture

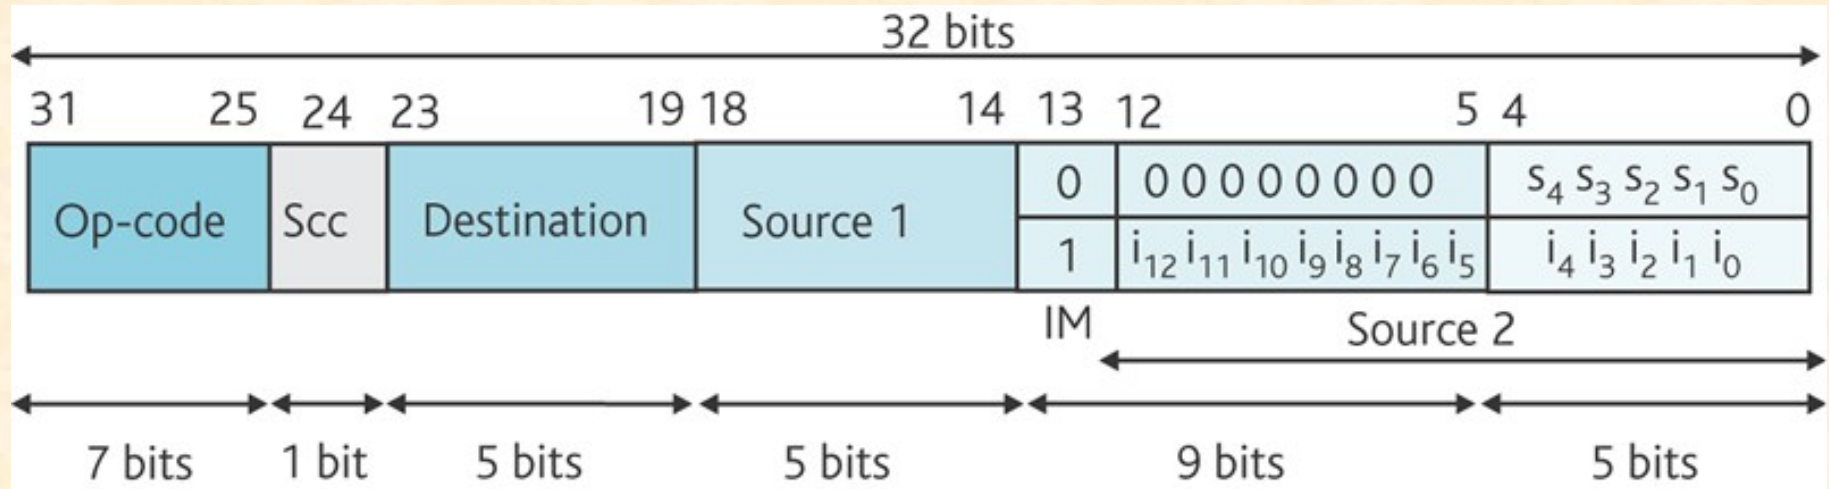
7. Because of (6) the RISC does not implement a microprogrammed architecture.
 - The distinction between machine cycle (number of clock cycles to complete an instruction) and microcycle (one clock cycle to complete a micro-operation) has vanished.
8. Single instruction format:
 - Decoding logic is simpler than for variable length instructions.
 - Memory usage may not be as efficient as variable length instructions.

The RISC Revolution

- Mid 1990s - Today
 - Distinction between RISC and CISC is blurred
 - Many RISC processors have become more complex than the CISC processors they were said to replace.
 - Many so called CISC processors have RISC like features.
 - Some RISC processors have more instructions than CISC processors.
 - RISC might be better referred to as Regular Instruction Set.

The Berkeley RISC: Instruction Format

(Led to the Commercial SPARC)



Scc: Specifies whether the condition code bits will be written to at the end of the instruction.

Oper: Each of five bit operands specifies one of 32 internal registers.

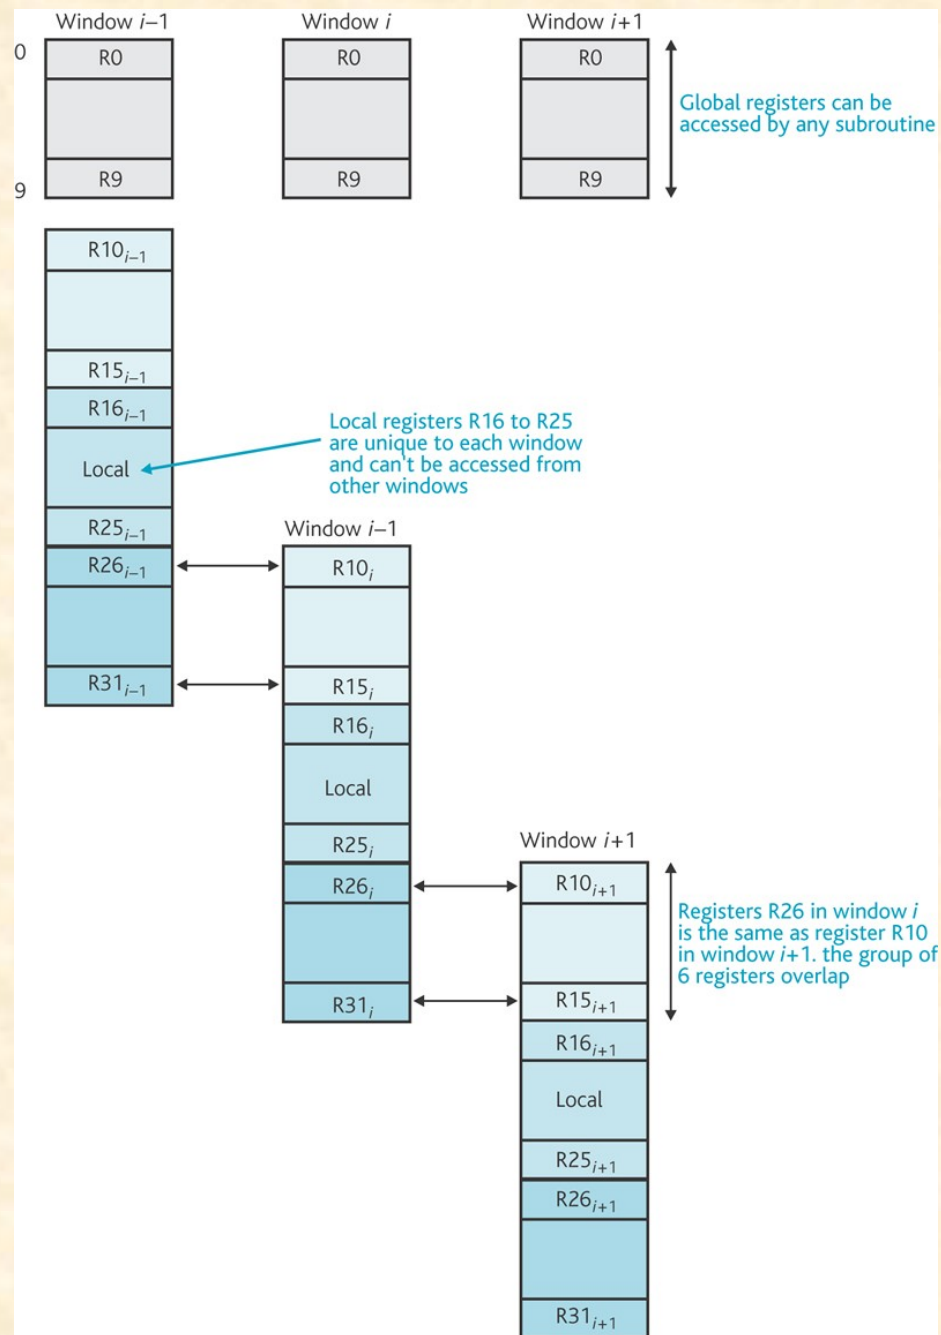
IM: If IM=0, [4..0] specify 2nd operand.

If IM=1, [12..0] specify a 13-bit constant, immediate value.

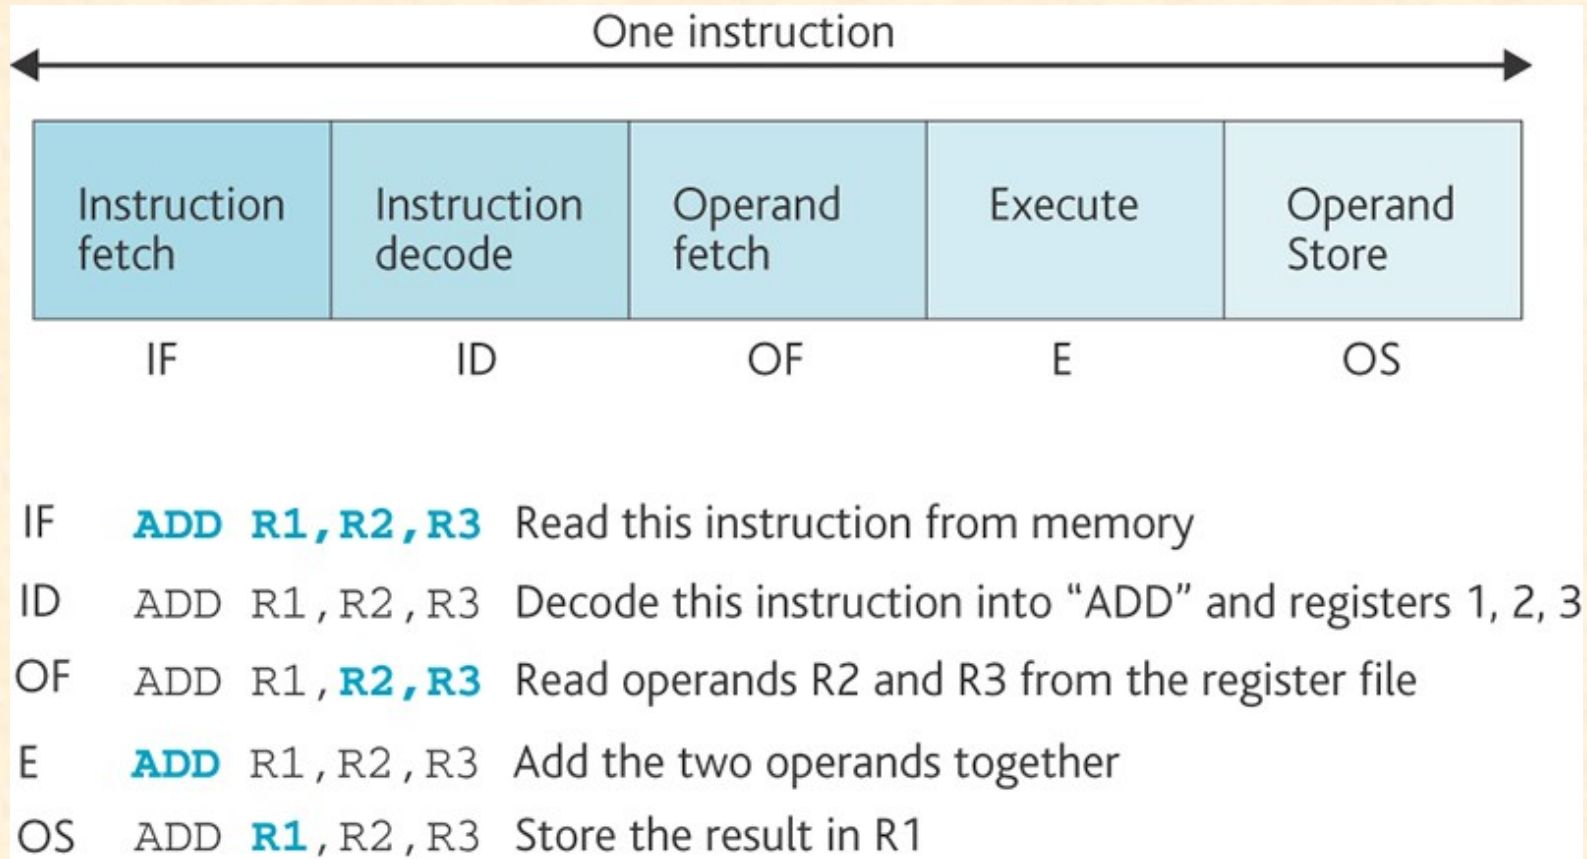
R1: Hardwired to 0. Constant and ADD R0,R1,R2 → MOVE R1,R2

The Berkeley RISC: Register Windows

- When a subroutine is called, the window pointer is incremented by 1.
- Program counter saved in R_d.
- The subroutine sees a different set of registers.
- 10 global + 8 x 10 local + 8 x 6 parameter transfer registers = 138 registers.
- Supports up to 8 nested subroutines.
- Main store is used if greater than 8.
- Context switching is expensive.

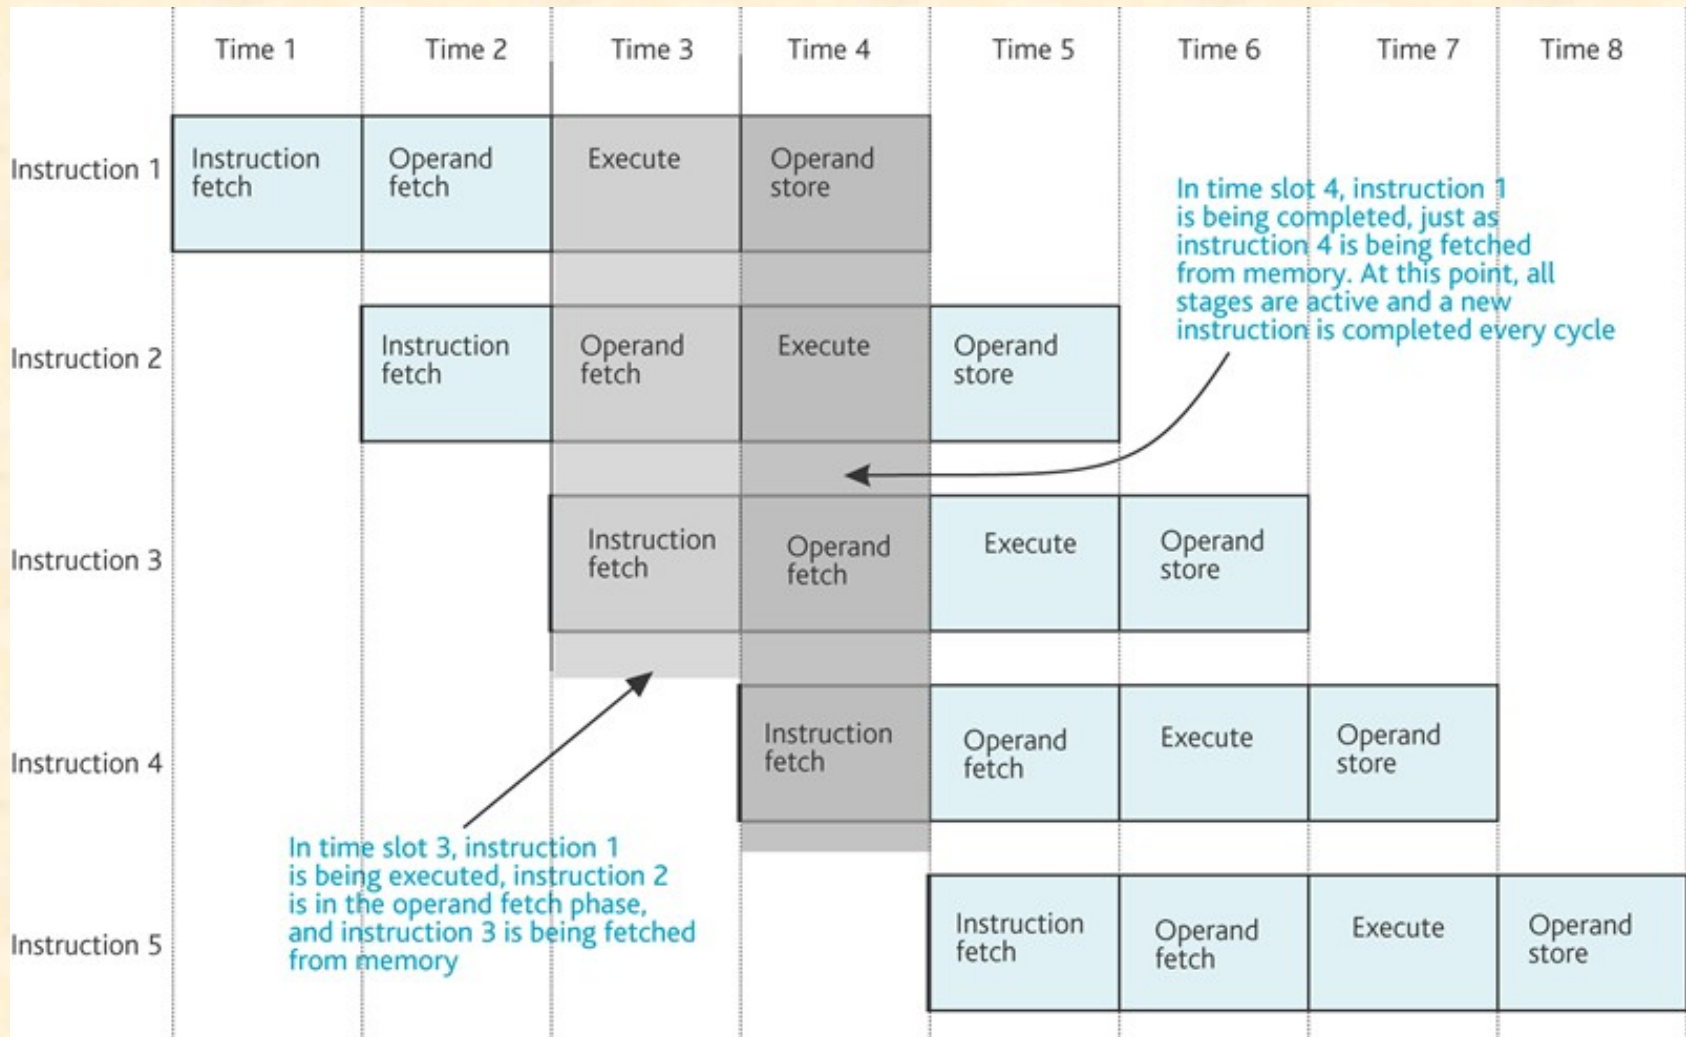


Instruction Execution Phases



- RISC processors don't need an Instruction Decode phase because their encoding is so simple.

Pipelining and Instruction Overlap



Pipelining Performance


- Consider the execution of n instructions using an m -stage pipeline.
- It will take m clock cycles for the 1st instruction to complete.
- The remaining $n - 1$ instructions execute at the rate of one clock cycle per instruction.
- The total time to execute the n instructions is:

$$m + (n - 1) \text{ cycles}$$

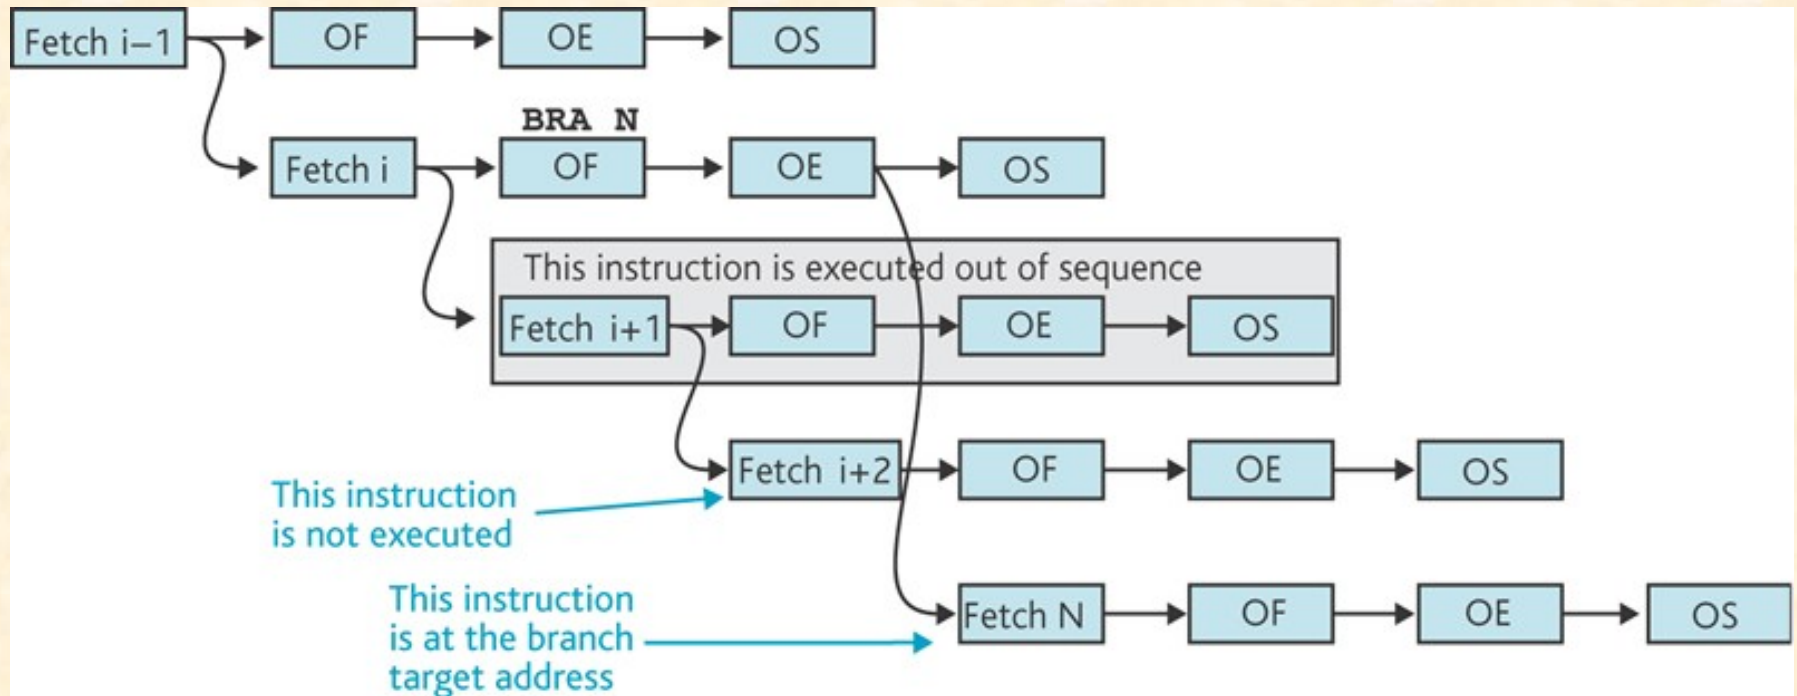
Pipelining Performance

Block Size	3-stage Pipeline	6-stage Pipeline	12-stage Pipeline
4	2.0000	2.6667	3.2000
8	2.4000	3.6923	5.0526
20	2.7272	4.8000	7.7419
100	2.9411	5.7143	10.8100
1000	2.9940	5.9701	11.8694
Infinite	3.0000	6.0000	12.0000

Pipeline Bubble



Overcoming The Pipeline Bubble



ADD R1,R2,R3

$[R3] \leftarrow [R1] + [R2]$

BRA N

GOTO ADDRESS N


ADD R2,R4,R5

$[R5] \leftarrow [R2] + [R4]$ This is executed.

ADD R7,R8,R9

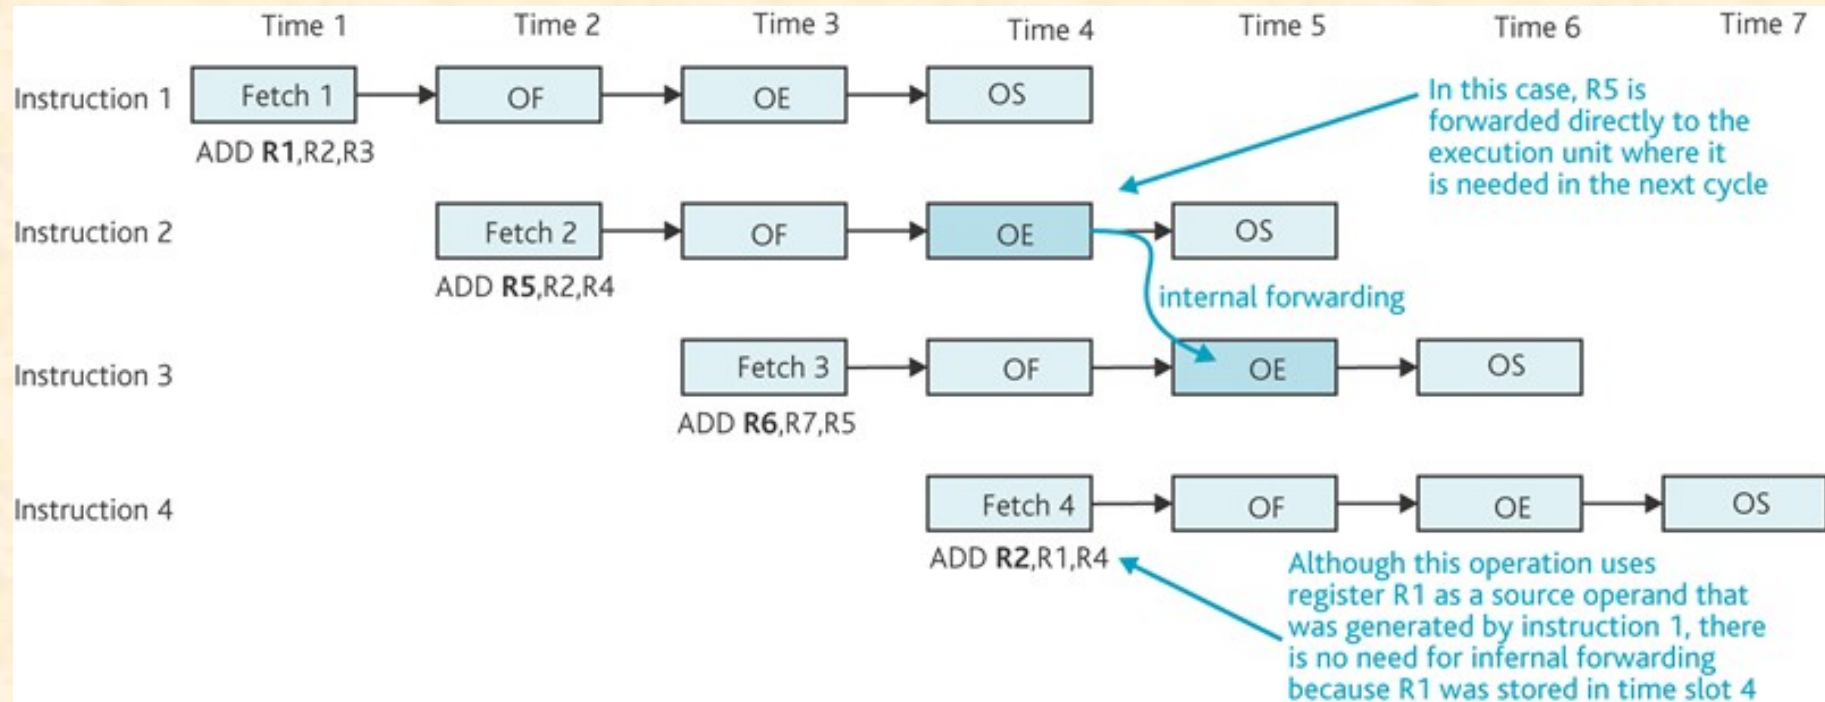
Not executed because branch is taken.

Data Dependency



- The pipeline is stalled after the fetch phase of instruction 3 for two clocked cycles.

Overcoming Data Dependency: Internal Forwarding




- | | | |
|----|----------------|-------------------------------|
| 1. | ADD R1, R2, R3 | $[R1] \leftarrow [R2] + [R3]$ |
| 2. | ADD R5, R2, R4 | $[R5] \leftarrow [R2] + [R4]$ |
| 3. | SUB R6, R7, R5 | $[R6] \leftarrow [R7] - [R5]$ |
| 4. | ADD R1, R1, R4 | $[R2] \leftarrow [R1] + [R4]$ |

Probabilistic Model of Instruction Execution

1. Each non-branch instruction is executed in one cycle.
 2. The probability that a given instruction is a branch is p_b .
 3. The probability that a branch instruction will be taken is p_t .
 4. If a branch is taken, the additional penalty is b cycles.
 5. If a branch is not taken, there is no penalty.
- The average time an instruction takes to execute is:

$$\begin{aligned} T_{ave} &= (1 - p_b) * 1 + p_b p_t (1 + b) + p_b (1 - p_t) * 1 \\ &= 1 + p_b p_t b \end{aligned}$$

Probabilistic Model of Branch Penalty



- The average number of cycles taken by a branch instruction is:

$$p_b (a (p_t p_c) + b (1 - p_t) (1 - p_c) + c p_t (1 - p_c) + d (1 - p_t) p_c)$$

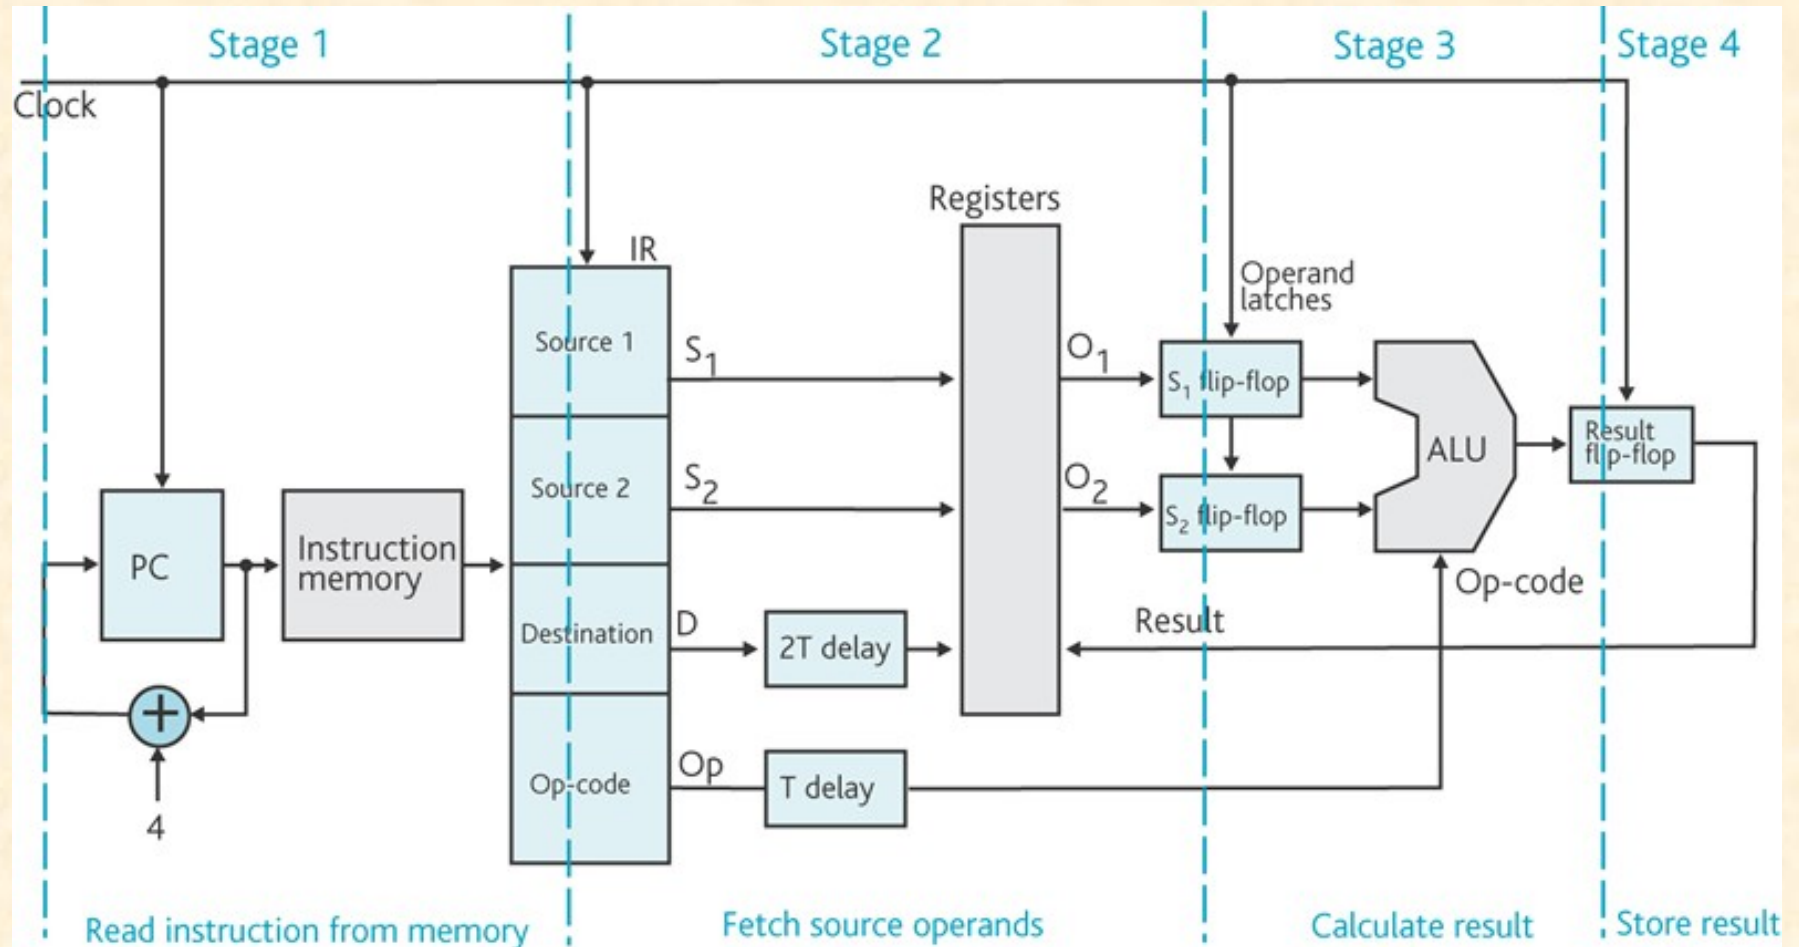
Implementing Branch Prediction

- Static Branch Prediction
 - Observation of real code has demonstrated $> 50\%$ chance that a branch will be taken
 - Fetch the next instruction at the target address.
- Some branch instructions are taken more or less frequently than others.
 - Basing the prediction on the opcode can yield as much as 75% accuracy.
 - Devote a bit in the opcode of the branch instruction
 - This bit is set if the compiler estimates a branch will be taken.
 - 75 – 94 % accuracy.

Implementing Branch Prediction

- Dynamic Branch Prediction
 - Prediction made at run time based on past behavior.
 - Processor uses a table that indicates the probability of each branch instruction.
 - The table is updated each time a branch instruction is executed.
 - Single bit branch prediction tables: 80%.
 - 5-bit branch prediction tables: 98%.

Using Latches to Implement Pipelining



Timing Diagram for a Pipelined Computer

